

Insurance auditing - decision log

Implementation candidate | Source commit 6fee1da60b74512156637a22be15d996a36627e1

Evidence and execution boundary

The Work-session assistant extracted and source-reviewed finite JSON rules and frozen service mappings. This acquisition was author review. Python replays saved records without an LLM/API. Accepted bundles bind source, schema, mapping and review identities. Bundle traces identify the controlling clause and partner context. Unknown operators and incompatible packages fail before pricing.

Service identity and incomplete opinions

Contract catalogs are not proof that an incomplete description identifies their only similar service. H1 development exposed one wrong specialty assignment (INV-H1-000236, 43,650 cents); all 39 analogous mapping keys were withdrawn. Essential qualifiers need independent description evidence; only generic noun elision retains a lower evidence grade. Billed rates/units never select service identity. If any necessary fact is unresolved, withhold the complete row; a known error alone does not establish the corrected total.

Quantities, history and allocation

Use supplied billed quantities under the contracts' pricing rules; a wrong unit label does not authorize a guessed conversion. Composite hour/item quantities lack a second dimension and remain unsupported. Prior usage is strictly ordered by service date and line ID, over the supplied term; quarantined headers retain patient candidates, and missing identity stays uncertain. Conflicting ownership and duplicate service/day allocation have no invented split. A supported single capped line can be limited to the contractual maximum. Round integer cents half-up after each adjustment.

Exclusions and amendments

H1/H2/H4 explicitly support bidirectional exclusions; H3/H5 direction is unresolved at nonzero distances. Use same-patient scope as the recorded contextual reading and withhold exactly-N-day boundary cases. H3 A1.1 applies seven repricings and two additions from service date 2025-01-01 despite the broad amendment header. A1.4.2 settlement protection uses the explicit interpretation that a valid invoice cannot settle before issue/service; no settlement field is fabricated.

Hospital-specific missing facts

H4 section 1.4 defines instance as unit, but section 8 leaves patient aggregation unclear: bound same-patient through all-patient usage and emit only invariant outcomes. H5 line facility is absent; adopt invoice facility from the supplied relational shape and section 10.1, as an explicit assumption. Outcome-relevant opinions are capped at .65 confidence. H2 supplies admission/discharge dates but lacks actual submission, detailed episode/leave and possible exception evidence; invoice date is not submission. All H2 full opinions are withheld. Its 07:00 Service Day is not established by calendar dates; day-dependent diagnostics remain bounded or uncertain.

Confidence, evaluation and remaining work

H1 patient groups were fixed before label development; the reserved check is now exposed and later checks are regressions. H1 sparse-error and target confidence are policy judgments, not demonstrated target calibration. Missing necessary facts cannot be repaired by a low score. Final mappings/rules, error history and omissions are retained. See [evaluation_report.md](#) and [docs/implementation_changes.md](#). The independent audit's two findings are corrected; closure re-audit and external delivery remain pending.